# Geometric Intelligence: A Differential-Geometric Decision-Making Framework on Data-Driven Differentiable Manifolds

Etale Cohomology

**Abstract**

We propose **Geometric Intelligence (GI) theory**, a mathematical framework that constructs corporate management environments and national policy environments as data-driven Riemannian manifolds using variational autoencoders (VAEs), and performs curvature tensor analysis, Lie derivative evaluation, geodesic calculation, and optimal control simulation on these manifolds to generate actionable intelligence for decision-making. Conventional data analysis—regression, PCA, linear optimization—rests on the implicit assumption that the environment is Euclidean. Real-world business and policy environments are inherently nonlinear: the same action produces qualitatively different outcomes depending on the agent's position in the state space. GI theory captures this nonlinear structure as curvature of a Riemannian manifold and applies the full differential-geometric toolkit—covariant derivatives, curvature tensors, Lie derivatives, and geodesics—to socioeconomic data. We establish a 10-step pipeline from data collection through manifold construction, metric learning, curvature computation, Lie derivative analysis, and optimal control simulation, demonstrated via two case studies. We further construct five pillars of reliability (MC Dropout, SHAP, DML, ZKML, formal verification), design an AI Agent architecture on a digital twin, and systematize seven geometric extensions from Morse theory to Sheaf theory. The central result (Proposition 2.1) shows that a VAE decoder satisfying four conditions—compactness, smoothness, full-rank Jacobian, and injectivity—yields an embedded submanifold with a well-defined Riemannian structure independent of industry or data distribution. We identify a four-layer structural gap explaining why this framework has not appeared in prior literature.

## 1 Introduction

### 1.1 Motivation: Navigating a Curved World with a Flat Map

Since Riemann's 1854 Habilitationsschrift "Über die Hypothesen, welche der Geometrie zu Grunde liegen," which founded the study of curved higher-dimensional spaces via the metric tensor, and Einstein's 1915 completion of general relativity—encoding gravity as curvature of a 4-dimensional pseudo-Riemannian manifold $(M, g)$ through the field equation

$$R_{\mu\nu} - \tfrac{1}{2} g_{\mu\nu} R + \Lambda g_{\mu\nu} = \frac{8\pi G}{c^4} T_{\mu\nu}, \tag{1}$$

physicists and engineers have routinely computed on manifolds. GPS satellite orbits require general-relativistic corrections (without which positional errors accumulate at $\sim$10 km/day), black hole imaging relies on tracing geodesics through curved spacetime, and robotic control exploits the Riemannian structure of configuration spaces such as the $n$-torus $T^n = (S^1)^n$ for $n$-joint manipulators.

In stark contrast, corporate and national-policy decision-making still relies on linear methods—regression, PCA, linear programming—which implicitly assume a Euclidean environment. This paper argues that this assumption is structurally inadequate and proposes Geometric Intelligence (GI) theory as an alternative.

### 1.2 Why Business and Policy Environments Are Curved

Three examples illustrate the inherent nonlinearity:

**Exchange-rate nonlinearity.** A 1% currency movement affects corporate earnings qualitatively differently at 90 JPY/USD (strong-yen regime) versus 150 JPY/USD (weak-yen regime). At the strong-yen extreme, an additional 1% appreciation can be catastrophic for exporters; at the weak-yen extreme, the marginal impact is limited. The linear model earnings $= a \times \text{FX} + b$ cannot capture this.

**Interest-rate nonlinearity.** A 25-basis-point rate hike triggers a qualitatively different market response at the zero lower bound (potential regime change) versus at 5% (marginal adjustment).

**Alliance-policy nonlinearity.** Strengthening an alliance acts as "insurance" during stable periods but may be perceived as "provocation" during tensions, potentially escalating the very instability it aims to counter.

In each case, the *same action* at *different positions* in the state space produces qualitatively different outcomes—the defining signature of a curved space in Riemannian geometry.

### 1.3 Contributions

1. **Manifold construction pipeline.** We construct data-driven Riemannian manifolds via the pullback metric $g_{ij}(z) = (J_{f_\theta}^\top J_{f_\theta})_{ij}$ of a VAE decoder and formalize the conditions guaranteeing manifold structure (Proposition 2.1).
2. **Differential-geometric analysis tools.** We apply scalar curvature maps, Lie derivatives, Lie brackets, geodesics, and Pontryagin's maximum principle to socioeconomic data.
3. **Five pillars of reliability.** MC Dropout, SHAP, DML, ZKML, and formal verification (Coq, Agda, Lean 4, Z3).
4. **AI Agent architecture.** Digital-twin manifold with deep RL agents; emergent phenomena analyzed via Lie bracket non-commutativity.
5. **Seven geometric extensions.** Morse theory, exotic manifolds, surgery theory, information geometry, symplectic geometry, Weyl geometry, and Sheaf theory.

### 1.4 Relation to Prior Work

GI theory sits at the intersection of four research areas.

**Latent space geometry.** Arvanitidis et al. [Arvanitidis et al., 2018] introduced the pullback metric to VAE latent spaces, demonstrating non-Euclidean geometry. GI theory extends this insight to the full chain: curvature tensors, Lie derivatives, and optimal control on socioeconomic data.

**Vector Diffusion Maps.** Singer & Wu [Singer & Wu, 2012] construct the connection Laplacian from discrete data with a spectral convergence guarantee. GI theory uses VDM as an independent validation of the VAE pullback metric.

**Neural ODE and optimal control.** Chen et al. [Chen et al., 2018] parameterize ODE right-hand sides with neural networks and compute gradients via the adjoint method—the discrete analog of Pontryagin's costate equation. GI theory applies Neural ODE to geodesic computation and optimal control on Riemannian manifolds.

**Geometric Deep Learning.** Bronstein et al. [Bronstein et al., 2021] present a unifying framework for geometric deep learning on grids, groups, graphs, geodesics, and gauges. However, their scope does not include constructing differentiable manifolds from socioeconomic data or performing differential-geometric analysis for decision-making. GI theory extends the geometric deep learning paradigm to this domain.

**Four-layer structural gap.** We identify why this framework has not appeared before: (1) data sparsity and educational gaps in social science; (2) the known-manifold assumption in engineering; (3) the two-manifold confusion in ML (NPG/TRPO/PPO address only the policy parameter manifold $\mathcal{M}_{\text{weights}}$, ignoring the environment state manifold $\mathcal{M}_{\text{env}}$); (4) the disciplinary intersection problem requiring simultaneous expertise in differential geometry, VAE/Neural ODE, deep RL, and domain-specific business data.

## 2 Mathematical Foundations

### 2.1 Topological Preliminaries

**Definition 2.1** (Topological space)**.** *A pair $(X, \mathcal{O})$ is a topological space if $\mathcal{O} \subset 2^X$ satisfies: (T1) $\emptyset, X \in \mathcal{O}$; (T2) arbitrary unions of elements of $\mathcal{O}$ belong to $\mathcal{O}$; (T3) finite intersections of elements of $\mathcal{O}$ belong to $\mathcal{O}$.*

**Definition 2.2** (Compactness)**.** *$X$ is compact if every open cover admits a finite subcover. By the Heine–Borel theorem, a subset of $\mathbb{R}^n$ is compact iff it is closed and bounded.*

**Definition 2.3** (Hausdorff property)**.** *$X$ is Hausdorff if any two distinct points admit disjoint open neighborhoods.*

### 2.2 Differentiable Manifolds

**Definition 2.4** (Chart and atlas)**.** *A $d$-dimensional chart on $M$ is a pair $(U, \varphi)$ where $U \subset M$ is open and $\varphi : U \to \varphi(U) \subset \mathbb{R}^d$ is a homeomorphism. A collection $\mathcal{A} = \{(U_\alpha, \varphi_\alpha)\}$ covering $M$ is a $C^k$ atlas if all transition maps $\varphi_\beta \circ \varphi_\alpha^{-1}$ are $C^k$.*

**Definition 2.5** ($C^k$ differentiable manifold)**.** *A second-countable Hausdorff space $M$ equipped with a maximal $C^k$ atlas is a $d$-dimensional $C^k$ differentiable manifold.*

**Smoothness requirement: $k \geq 3$.** The curvature tensor ${R^l}_{ijk}$ requires second derivatives of the metric. Since $g_{ij} \in C^{k-1}$ for $f_\theta \in C^k$, we need $k - 1 \geq 2$, i.e., $k \geq 3$. Using $C^\infty$ activations ($\tanh$, GELU, Softplus) guarantees this. ReLU $= \max(0, x)$ is $C^0 \setminus C^1$ at $x = 0$ and is therefore excluded.

### 2.3 Immersions, Embeddings, and the VAE Manifold Theorem

**Definition 2.6** (Immersion). *A $C^k$ map $f : M \to N$ is an immersion if $df_p : T_pM \to T_{f(p)}N$ is injective at every $p \in M$. Equivalently,* $\operatorname{rank} J_f(p) = \dim M$ *for all $p$.*

**Definition 2.7** (Embedding). *An immersion $f : M \to N$ is an embedding if it is injective and $f : M \to f(M)$ is a homeomorphism (with the subspace topology on $f(M)$).*

**Theorem 2.1** (Compact immersion $\Rightarrow$ embedding; Lee [Lee, 2012, Thm. 4.25]). *If $M$ is compact Hausdorff, $N$ is Hausdorff, and $f : M \to N$ is a continuous injective immersion, then $f$ is an embedding.*

*Proof sketch.* It suffices to show $f^{-1} : f(M) \to M$ is continuous. For any closed $C \subset M$: $M$ compact and $C$ closed $\Rightarrow C$ compact $\Rightarrow f(C)$ compact (continuous image) $\Rightarrow f(C)$ closed ($N$ Hausdorff). Thus $f$ maps closed sets to closed sets, so $f^{-1}$ is continuous. □

**Proposition 2.1** (VAE decoder manifold structure). *Let $f_\theta : U \to \mathbb{R}^n$ be a VAE decoder satisfying:*

*(i) $U \subset \mathbb{R}^d$ is compact (closed and bounded);*

*(ii) $f_\theta \in C^k$, $k \geq 3$ (guaranteed by $C^\infty$ activations; ReLU excluded);*

*(iii)* $\operatorname{rank} J_{f_\theta}(z) = d$ *for all $z \in U$ (verified via SVD: $\sigma_{\min}(J) > 0$);*

*(iv) $f_\theta|_U$ is injective.*

*Then $f_\theta(U)$ is a $d$-dimensional $C^k$ embedded submanifold of $\mathbb{R}^n$, and*

$$g_{ij}(z) = \sum_{a=1}^{n} \frac{\partial f_\theta^a}{\partial z^i} \frac{\partial f_\theta^a}{\partial z^j} = (J_{f_\theta}^\top J_{f_\theta})_{ij} \tag{2}$$

*is a Riemannian metric on $U$.*

*Proof.* **Step 1 (Immersion).** $f_\theta \in C^k$ and $\operatorname{rank} J_{f_\theta}(z) = d$ at every $z \in U$ imply that $f_\theta|_U$ is a $C^k$ immersion.

**Step 2 (Embedding).** $U$ is compact (hypothesis (i)), $\mathbb{R}^n$ is Hausdorff, and $f_\theta|_U$ is a continuous injective immersion. By Theorem 2.1, $f_\theta|_U$ is a $C^k$ embedding, so $f_\theta(U)$ is a $C^k$ embedded submanifold (Lee [Lee, 2012, Prop. 5.2]).

**Step 3 (Positive definiteness).** For any $0 \neq v \in \mathbb{R}^d$, full rank implies $J_{f_\theta} v \neq 0$, hence $v^\top g(z) v = \|J_{f_\theta} v\|^2 > 0$. Symmetry is immediate from $(J^\top J)_{ij} = (J^\top J)_{ji}$. □

**Remark 2.1** (Industry independence). *Positive definiteness of $g_{ij}$ follows from the algebraic structure $J^\top J$ alone, independent of the industry, data distribution, or statistical properties. This guarantees Riemannian structure uniformly across all GI theory applications.*

**Failure modes.** (i) $U$ non-compact $\Rightarrow$ embedding upgrade fails. (ii) ReLU $\Rightarrow$ immersion undefined at fold hyperplanes; second derivatives needed for curvature are undefined. (iii) Rank deficiency $\Rightarrow g_{ij}$ acquires zero eigenvalues; $g^{ij}$ does not exist; Christoffel symbols undefined. (iv) Non-injective $\Rightarrow$ self-intersection; image is not a manifold.

## 3 Tangent Spaces and the Riemannian Metric

**Definition 3.1** (Tangent vector). *For a smooth curve $\gamma : (-\varepsilon, \varepsilon) \to M$ with $\gamma(0) = p$ and chart $(U, \varphi)$, the tangent vector at $p$ is $v = \frac{d}{dt}\big|_{t=0} (\varphi \circ \gamma)(t) \in \mathbb{R}^d$. The tangent space $T_pM$ is a $d$-dimensional vector space.*

**Definition 3.2** (Riemannian metric). *A Riemannian metric on $M$ is a smoothly varying family of positive-definite symmetric bilinear forms $g_p : T_pM \times T_pM \to \mathbb{R}$. In local coordinates: $ds^2 = g_{ij}(z)\, dz^i\, dz^j$.*

**Definition 3.3** (Pullback metric). *For a VAE decoder $f_\theta : U \subset \mathbb{R}^d \to \mathbb{R}^n$ satisfying Proposition 2.1:*

$$g_{ij}(z) = (J_{f_\theta}^\top J_{f_\theta})_{ij}, \qquad J_{f_\theta}(z) = \left( \frac{\partial f_\theta^a}{\partial z^i} \right) \in \mathbb{R}^{n \times d}. \tag{3}$$

*The eigenvalues of $g_{ij}(z)$ quantify directional sensitivity: large eigenvalues correspond to directions in which small latent movements cause large data-space changes.*

**Definition 3.4** (Curve length and energy)**.** *For a smooth curve $\gamma : [a, b] \to M$:*

$$L[\gamma] = \int_a^b \sqrt{g_{ij}(\gamma)\, \dot{\gamma}^i \dot{\gamma}^j}\, dt, \tag{4}$$

$$E[\gamma] = \tfrac{1}{2} \int_a^b g_{ij}(\gamma)\, \dot{\gamma}^i \dot{\gamma}^j\, dt. \tag{5}$$

*Critical points of $E$ with fixed endpoints satisfy the geodesic equation* (14)*.*

# 4 Connection, Covariant Derivative, and Curvature

**Theorem 4.1** (Levi-Civita connection)**.** *On a Riemannian manifold $(M, g)$, there exists a unique affine connection $\nabla$ that is (LC1) torsion-free ($\Gamma^k_{ij} = \Gamma^k_{ji}$) and (LC2) metric-compatible ($\nabla_k g_{ij} = 0$). Its Christoffel symbols are:*

$$\Gamma^k_{ij}(z) = \tfrac{1}{2} g^{kl}(z)(\partial_i g_{jl} + \partial_j g_{il} - \partial_l g_{ij})\,. \tag{6}$$

The covariant derivative of a vector field $W^k$ along $\partial_i$ is:

$$\nabla_i W^k = \partial_i W^k + \Gamma^k_{ij} W^j. \tag{7}$$

The first term captures component change; the second corrects for coordinate curvature.

**Definition 4.1** (Riemann curvature tensor)**.**

$$R^l{}_{ijk} = \partial_j \Gamma^l_{ik} - \partial_i \Gamma^l_{jk} + \Gamma^l_{jm}\Gamma^m_{ik} - \Gamma^l_{im}\Gamma^m_{jk}. \tag{8}$$

*The Ricci tensor is* $\mathrm{Ric}_{ij} = R^l{}_{ilj}$ *and the scalar curvature is* $\mathrm{Scal} = g^{ij}\mathrm{Ric}_{ij}$.

**Interpretation.** $\mathrm{Scal}(z) > 0$: geodesic ball volume $<$ Euclidean (converging geodesics, stable structure). $\mathrm{Scal}(z) < 0$: volume $>$ Euclidean (diverging geodesics, unstable structure). This information is *principally inaccessible* to PCA, which projects onto linear subspaces and cannot define curvature.

**Autodiff full-chain computation.** Using PyTorch `autograd` with `create_graph=True`: (1) Jacobian $J_{f_\theta}$ $\to$ pullback metric $g_{ij}$; (2) $\partial_i g_{jl} \to$ Christoffel symbols $\Gamma^k_{ij}$; (3) $\partial_j \Gamma^l_{ik} \to$ curvature tensor $R^l{}_{ijk}$; (4) vector field derivatives $\to$ Lie brackets and Lie derivatives.

# 5 Lie Derivatives and Killing Fields

**Definition 5.1** (Lie derivative of the metric)**.** *For a vector field $V$:*

$$(\mathcal{L}_V g)_{ij} = \nabla_i V_j + \nabla_j V_i. \tag{9}$$

*If $\mathcal{L}_V g = 0$ (Killing equation), $V$ is a Killing field generating isometries.*

In GI theory, policies and managerial actions are modeled as vector fields. $\mathcal{L}_V g = 0$ means the action *preserves the environmental structure* while moving the state—a gentle policy. $\mathcal{L}_V g \neq 0$ means the action *deforms the structure itself*: which states are nearby, which are distant, which are stable, which are unstable. This is qualitatively different from conventional policy evaluation (DID, regression discontinuity), which measures *outcome differences*, not *geometric deformation*.

**Definition 5.2** (Lie bracket)**.** *For vector fields $V_A, V_B$:*

$$[V_A, V_B]^k = V^i_A \partial_i V^k_B - V^i_B \partial_i V^k_A. \tag{10}$$

*$[V_A, V_B] = 0 \Rightarrow$ the two policies commute (order-independent). $[V_A, V_B] \neq 0 \Rightarrow$ the outcome depends on execution order. $\|[V_A, V_B]\|$ measures interference strength; the direction indicates the interference vector.*

# 6 Geometric AI Toolkit

## 6.1 VAE with Smooth Decoder

The VAE decoder $f_\theta : \mathbb{R}^d \to \mathbb{R}^n$ is a continuous function that generates data-space points for *any* latent point, filling the space smoothly between training points. This is the critical difference from UMAP/t-SNE, which produce *discrete point configurations* where (i) no Riemannian metric is defined, (ii) no derivatives exist, (iii) no tensor fields can be defined, (iv) no control-system formulation is possible. VAE creates a smooth manifold; UMAP creates a scatter plot.

### 6.2 Vector Diffusion Maps (VDM)

Singer & Wu [Singer & Wu, 2012] construct the graph connection Laplacian:

$$(S_\varepsilon v)_i = \frac{1}{d_i}\sum_j w_{ij}(v_i - O_{ij}v_j), \quad w_{ij} = e^{-\|z_i - z_j\|^2/\varepsilon}, \quad d_i = \textstyle\sum_j w_{ij}, \tag{11}$$

where $O_{ij} \in O(d)$ are optimal rotation matrices between local tangent spaces.

**Theorem 6.1** (Spectral convergence; Singer & Wu, 2012)**.** *Under Assumption A1, as $N \to \infty$, $\varepsilon \to 0$ with $N\varepsilon^{d/2+1} \to \infty$, $S_\varepsilon$ converges spectrally to the connection Laplacian on the tangent bundle of $M_0$.*

### 6.3 Neural ODE and Optimal Control

Neural ODE [Chen et al., 2018]: $\dot{z} = f_\theta(z,t)$, with gradients computed via the adjoint method:

$$\frac{da}{dt} = -a(t)^\top \frac{\partial f_\theta}{\partial z}(z(t),t), \tag{12}$$

where $a(t) = \partial L/\partial z(t)$ is the adjoint state—the discrete analog of Pontryagin's costate equation.

For an affine control system $\dot{z}^k = F^k(z) + u_\alpha B^k_\alpha(z)$, Pontryagin's maximum principle gives:

$$H(z,p,u) = p_k\big(F^k + u_\alpha B^k_\alpha\big) - \ell(z,u), \qquad \dot{z}^k = \frac{\partial H}{\partial p_k}, \quad \dot{p}_k = -\frac{\partial H}{\partial z^k}. \tag{13}$$

The geodesic equation is recovered when $u = 0$ and $F^k = 0$:

$$\ddot{\gamma}^k + \Gamma^k_{ij}\dot{\gamma}^i\dot{\gamma}^j = 0. \tag{14}$$

## 7 The 10-Step Pipeline

1. **Data collection.** Industry-specific indicators.
2. **VAE manifold construction.** Train VAE with $C^\infty$ decoder; verify Proposition 2.1 conditions.
3. **Pullback metric.** $g_{ij}(z) = (J^\top J)_{ij}$ via autodiff.
4. **Geometric structure extraction.** VDM for independent validation.
5. **Christoffel symbols.** Eq. (6) via autodiff.
6. **Curvature tensor and scalar curvature.** Eq. (8). Map stable ($\mathrm{Scal} > 0$) vs. unstable ($\mathrm{Scal} < 0$) regions.
7. **Lie derivative.** Eq. (9). Lie bracket analysis via Eq. (10).
8. **Geodesics.** Eq. (14) via Runge–Kutta.
9. **Optimal control.** Pontryagin's maximum principle, Eq. (13).
10. **Visualization and decision-making.** 3D terrain maps and reports.

## 8 Assumptions and Data Quality

**Assumption 1** (Manifold hypothesis (A1))**.** *Data $X = \{x_i\} \subset \mathbb{R}^n$ are noisy samples from a neighborhood of a $d$-dimensional compact connected $C^\infty$ submanifold $M_0$.*

**Assumption 2** (Decoder regularity (A2))**.** *(a) $f_\theta \in C^k$, $k \geq 3$. (b) $\operatorname{rank} J_{f_\theta}(z) = d$ for all $z \in U$.*

**Assumption 3** (Sufficient data (A3))**.** *$N\varepsilon^{d/2+1} \to \infty$. Rule of thumb: $N > 100 \times 2^d$.*

**Assumption 4** (Numerical stability (A4))**.** *$\kappa(g(z))$ is bounded. Regions with $\kappa > 10^4$ are flagged as numerically unreliable.*

**Five data quality checks.** (1) PCA scree plot: $\geq 80\%$ variance in few components. (2) Data density: $N > 100 \times 2^d$. (3) Reconstruction error $< 10\%$; $\sigma_{\min}(J) > 10^{-4}$. (4) Robustness to seed/bandwidth changes. (5) Curvature patterns consistent with known historical events.

## 9 Five Pillars of Reliability

### 9.1 Pillar 1: MC Dropout for Uncertainty Quantification

MC Dropout [Gal & Ghahramani, 2016] keeps dropout active at inference time, running $T$ forward passes to obtain $\hat{y}_1, \ldots, \hat{y}_T$. The mean $\bar{y} = \frac{1}{T}\sum_t \hat{y}_t$ is the point estimate; the variance $\sigma^2 = \frac{1}{T}\sum_t (\hat{y}_t - \bar{y})^2$ quantifies epistemic uncertainty. Gal & Ghahramani showed this corresponds to variational inference in a Bayesian neural network.

### 9.2 Pillar 2: SHAP for Explainability

SHAP (SHapley Additive exPlanations; Lundberg & Lee 2017), based on Shapley values [Shapley, 1953], decomposes predictions into per-feature contributions. Spatial SHAP on the manifold reveals how factor importance varies with position.

### 9.3 Pillar 3: DML for Causal Inference

Double Machine Learning [Chernozhukov et al., 2018] separates confounding from treatment effects via cross-fitting. In Case Study A, SHAP ranks "advertising spend" first, but DML reveals "R&D investment" has the largest *causal* effect—a priority reversal caused by seasonal confounding.

### 9.4 Pillar 4: ZKML for Cryptographic Verification

ZKML applies zero-knowledge proofs (Goldwasser, Micali & Rackoff, 1985) to ML computations: the prover generates a succinct proof ($O(1)$ bytes) that the agreed-upon model architecture was executed correctly, verifiable without seeing data or parameters. Limitation: ZKML guarantees *computational correctness*, not *parameter legitimacy*. Three defense layers are needed: (1) ZKML, (2) training-log commitment (hash chain), (3) third-party audit.

### 9.5 Pillar 5: Formal Verification

Theorem provers (Coq, Agda, Lean 4) and SMT solvers (Z3) verify logical consistency of the entire pipeline.

## 10 AI Agent Architecture on Digital Twins

The management-environment manifold is instantiated as a server-side digital twin. Fourteen AI Agents, trained via deep reinforcement learning, operate on this manifold: executing geodesic-following movements, computing Lie derivatives of proposed actions, and running optimal-control simulations. Emergent phenomena—whole-system behaviors inexplicable by summing individual actions—are analyzed via Lie bracket non-commutativity: $[V_A, V_B] \neq 0$ implies execution-order dependence and unpredictable interaction. A three-level monitoring system prevents runaway: (1) parameter bounds, (2) curvature-based anomaly detection, (3) formal verification of logical consistency.

## 11 Geometric Extensions

**Morse theory.** Critical points of the scalar curvature function classify tipping points; Morse inequalities bound their count. **Exotic manifolds & optimal transport.** Homeomorphic but non-diffeomorphic structures model environments with identical topology but distinct smoothness—applicable to early pandemic detection. **Surgery theory.** Controlled topological modification of manifolds models radical restructuring (M&A, regime change). **Symplectic geometry.** Symplectic integrators preserve Hamiltonian structure for long-horizon simulation stability. **Weyl geometry.** Gauge invariance ensures analysis results are independent of currency denomination—the sign of scalar curvature is gauge-invariant. **Sheaf theory.** The Sheaf Laplacian detects local inconsistencies in data or model outputs, flagging tampering or assumption violations.

## 12 Case Studies

### 12.1 Case Study A: Manufacturing Company X

A mid-cap manufacturer (revenue ~¥100B) considering emerging-market expansion. Dataset: 30 indicators $\times$ 6,000 data points $\rightarrow$ 6-dimensional manifold.

**Proposition 2.1 verification:** (i) $\|z\| \leq 3$ is compact ✓. (ii) $\tanh$ activations ✓. (iii) $\sigma_{\min} = 8.3 \times 10^{-3} > 10^{-4}$ ✓. (iv) No collisions in 1,000 random pairs ✓.

**Findings:** Negative scalar curvature region identified on the emerging-market entry path (structurally unstable zone). Geodesic-based optimal path circumvents this region. Lie derivative analysis: cost-reduction policy is structure-preserving ($\|\mathcal{L}_V g\| \approx 0$); market-expansion policy is structure-deforming. Lie bracket: cost-reduction and market-expansion interfere ($[V_{\text{cost}}, V_{\text{market}}] \neq 0$).

SHAP ranks advertising first; DML reveals R&D investment has the largest causal effect—priority reversal due to seasonal confounding.

### 12.2 Case Study B: National Security Council

193 countries $\times$ 40 indicators $\times$ 30 years $\rightarrow$ 8-dimensional manifold. Indicators: military spending/GDP, nuclear forces, alliance networks, UN voting patterns, trade dependency matrices, energy import concentration, democracy indices.

**Data-sparse regions:** For authoritarian states with limited data disclosure, $\sigma_{\min}$ drops to $3.2 \times 10^{-5}$ (below threshold). Analysis for G20 + major allies proceeds; data-sparse regions flagged with "medium confidence"; additional intelligence collection recommended for the sparsest regions.

Alliance-strengthening policy CATE (DML) is strongly position-dependent on the manifold: large in weakly-allied regions, small where alliances are already strong (diminishing marginal returns).

## 13 Ethics and Limitations

**Conflicts of interest.** Simultaneous contracts with competing firms or governments create ethical risks. ZKML's three defense layers partially mitigate this, but data-selection fairness requires institutional safeguards (third-party audit).

**Incompleteness theorem of geometric modeling.** When Proposition 2.1 conditions or Assumptions A1–A4 fail, results are unreliable. Explicitly delineating the boundary between tractable and intractable problems is itself a contribution.

**Information symmetry/asymmetry.** Game-theoretic analysis of scenarios where adversaries operate equivalent AI systems.

## 14 Conclusion

We have proposed Geometric Intelligence (GI) theory: a framework that constructs data-driven Riemannian manifolds via VAE pullback metrics and applies the full differential-geometric toolkit—covariant derivatives, curvature tensors, Lie derivatives, geodesics, and optimal control—to corporate and national-policy decision-making.

The central claim is that business environments are *curved, not flat*: ignoring this curvature induces the same structural error as navigating the Earth's surface with a flat map. GI theory corrects this error with mathematical rigor.

Five pillars of reliability ensure trustworthy outputs. An AI Agent architecture enables autonomous policy exploration on the manifold. Seven geometric extensions expand the theoretical reach. The framework's limits—dependent on Proposition 2.1's four conditions and Assumptions A1–A4—are explicitly delineated.

Five guiding principles govern the pipeline: (1) *Purpose first*: generate intelligence one can act on. (2) *Don't add; raise reliability*. (3) *Separate the known from the unknown*. (4) *Data quality is everything*. (5) *Mathematics does not bend*: scalar curvature sign is the same regardless of who computes it.

## References


Amari, S.-I. (1998). Natural gradient works efficiently in learning. *Neural Computation*, 10(2), 251–276.

Arvanitidis, G., Hansen, L.K., & Hauberg, S. (2018). Latent space oddity: On the curvature of deep generative models. *ICLR*.

Bronstein, M.M., Bruna, J., LeCun, Y., Szlam, A., & Vandergheynst, P. (2021). Geometric deep learning: Grids, groups, graphs, geodesics, and gauges. *arXiv:2104.13478*.

Chen, R.T.Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2018). Neural ordinary differential equations. *NeurIPS*.

Chernozhukov, V., et al. (2018). Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 21(1), C1–C68.

do Carmo, M.P. (1992). *Riemannian Geometry*. Birkhäuser.

Etale Cohomology. (2024a). Geometric Intelligence, Volume 1. Zenodo. https://doi.org/10.5281/zenodo.19140918.

Etale Cohomology. (2024b). Geometric Intelligence, Volume 2. Zenodo. https://doi.org/10.5281/zenodo.19157891.

Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *ICML*.

Goldwasser, S., Micali, S., & Rackoff, C. (1985). The knowledge complexity of interactive proof systems. *STOC*.

Lee, J.M. (2012). *Introduction to Smooth Manifolds*. 2nd ed. Springer.

Lundberg, S.M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *NeurIPS*.

Milnor, J.W. (1963). *Morse Theory*. Princeton University Press.

Pontryagin, L.S., et al. (1962). *The Mathematical Theory of Optimal Processes*. Wiley.

Shapley, L.S. (1953). A value for $n$-person games. *Contributions to the Theory of Games*, 2, 307–317.

Singer, A., & Wu, H.-T. (2012). Vector diffusion maps and the connection Laplacian. *Comm. Pure Appl. Math.*, 65(8), 1067–1144.

Srivastava, N., et al. (2014). Dropout: A simple way to prevent neural networks from overfitting. *JMLR*, 15, 1929–1958.

Kingma, D.P., & Welling, M. (2014). Auto-encoding variational Bayes. *ICLR*.

Hafner, D., Pasukonis, J., Ba, J., & Lillicrap, T. (2023). Mastering diverse domains with world models. *arXiv:2301.04104*.

Kakade, S.M. (2001). A natural policy gradient. *NeurIPS*, 14.

Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv:1707.06347*.

McInnes, L., Healy, J., & Melville, J. (2018). UMAP: Uniform Manifold Approximation and Projection. *arXiv:1802.03426*.

van der Maaten, L., & Hinton, G. (2008). Visualizing data using t-SNE. *JMLR*, 9, 2579–2605.

Facco, E., et al. (2017). Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7, 12140.